

# AI Health Data Engine - Final PRD

## Product Overview

I developed an AI-powered health data extraction tool to address a critical problem in patient engagement with digital health applications. Traditional health apps require patients to navigate rigid data entry interfaces - dropdown menus, sliders, and structured forms - creating substantial cognitive burden that leads to poor engagement and incomplete datasets. This friction undermines both clinical decision-making and medical research effectiveness.

My solution transforms unstructured patient narratives into structured, standardized data using Generative AI. Rather than forcing patients to categorize symptoms on predetermined scales, they can input natural language descriptions of their health experiences. The system processes these narratives through Google's Gemini Flash API and outputs clean JSON data immediately suitable for clinical analysis.

The current implementation demonstrates a complete React frontend with voice input capabilities, real-time AI processing, and comprehensive data visualization. All core functionality is operational: natural language input, AI extraction, and structured output generation. The primary limitation is the absence of backend infrastructure, which was a deliberate choice to focus development effort on validating the AI extraction concept.

Target users include healthcare application developers who need reliable, complete patient-generated health data, and patients managing chronic conditions who require low-friction methods for daily health logging. The implementation validates both technical feasibility and user experience potential of AI-mediated health data extraction.

## Implementation Status and Core Features

The prototype includes fully implemented natural language input interfaces supporting both text entry and voice input via the browser's SpeechRecognition API. The AI-powered extraction pipeline processes narratives through Gemini Flash and returns structured data in real-time. The system organizes extracted information into clinically relevant categories including symptoms, meals, medications, stress levels, sleep patterns, and other health indicators.

Additional implemented features include example narrative generation to guide effective user input, comprehensive summary displays formatted for clinical review, and toggle functionality between human-readable summaries and raw JSON output. The entire user flow from narrative input to structured data presentation operates seamlessly.

The notable omission is backend infrastructure. This decision enabled rapid prototyping and concept validation while acknowledging that production deployment would require secure, local processing to protect patient data. The current architecture demonstrates proof-of-concept effectiveness while clearly flagging necessary changes for clinical implementation.

## AI Integration and Technical Architecture

The AI component performs zero-shot information extraction and semantic transformation on patient health narratives. Google Gemini Flash processes unstructured text inputs to identify and extract specific clinical data points while performing semantic mapping that converts natural language descriptions like "feeling bloated" into appropriate severity scores on clinical scales.

The extraction methodology derives directly from my thesis research, where I developed and validated tool calling approaches for maximizing accuracy in clinical data extraction. This established prompting infrastructure ensures consistent extraction quality without requiring rebuilding of validated methodologies.

The technical implementation uses React with TypeScript for component-based frontend architecture, managed through React hooks for state handling. Styling utilizes Tailwind CSS to maintain a clean, clinical aesthetic. AI integration operates through direct API calls to Google AI Studio with proper error handling and asynchronous processing capabilities.

The user interaction flow follows a straightforward pattern: patients input narratives via text or voice, the frontend transmits data to Gemini Flash API, AI processing extracts and structures information, and JSON responses populate both summary displays and raw data views. The system includes appropriate loading indicators and error handling throughout the processing pipeline.

Critical architectural decisions include frontend-only implementation for rapid prototyping, stateless processing with no data persistence, and direct API integration without backend abstraction. These choices enabled concept validation while creating privacy limitations that would require resolution before clinical deployment.

## Development Experience with AI Assistance

Building this project provided valuable insights into effective AI-assisted development practices. I relied primarily on Google AI Studio and Gemini for both product functionality and development assistance, though the relationship between AI help and human oversight proved more nuanced than anticipated.

AI assistance proved most effective for discrete, well-defined implementation tasks. Specific prompts such as "generate a React component for health data summary display with categorical sections" or "implement voice input toggle functionality with proper cleanup and error handling" produced solid foundational code requiring minimal modification. The AI excelled at generating

component structures, implementing browser APIs, and creating TypeScript interface definitions.

Where AI assistance was less effective involved architectural decisions, user experience flow, and integration between components. Requests involving multiple simultaneous changes often broke existing functionality while implementing new features. This experience led me to adopt a more methodical approach: using AI for specific implementation tasks while personally handling integration work to ensure system coherence.

The debugging experience highlighted similar patterns. AI could identify syntax errors and suggest fixes for isolated code blocks effectively, but struggled with complex state management issues or problems spanning multiple components. All user experience and design decisions remained entirely human-driven, including the clinical aesthetic, information hierarchy, and workflow design.

Prompt evolution throughout development demonstrated that broad initial requests produced generic, unusable results, while specific, task-oriented prompts yielded practical code. This required developing skills in decomposing complex requirements into discrete, actionable AI tasks while maintaining human control over system architecture and user experience design.

## Current Limitations and User Experience

The intended user experience centers on frictionless health data entry through a clean, clinical interface with clear value proposition messaging. Users can input health narratives via text entry or voice activation, guided by example content demonstrating effective input patterns. Real-time processing indicators provide feedback during AI extraction and structuring.

Results presentation includes categorized health summaries formatted for clinical review and toggle options for raw data comparisons between original narratives and structured JSON output. Navigation enables users to easily input additional entries while maintaining clear workflow progression.

Several significant limitations affect current implementation. Most critically, patient narratives transmit to Google's external API infrastructure, creating privacy concerns unacceptable for clinical deployment. While data contains no explicit patient identifiers, health narratives inherently include sensitive personal information.

The system provides no data persistence, making each session isolated without storage capabilities or user accounts. Accuracy validation has not been conducted specifically for this implementation, though extraction methodology derives from validated thesis research. Generated example narratives may reflect model biases that could influence authentic patient reporting patterns.

Additional limitations include browser compatibility requirements for voice input functionality and potential over-reliance risks where users might trust AI extraction accuracy without appropriate

clinical validation. The system includes clear disclaimers about non-medical applications, but appropriate trust boundaries require ongoing attention.

## Future Development Direction

Immediate technical improvements focus on privacy and security through implementation of secure backend infrastructure to eliminate external API data transmission. Adding local data storage and user session management would enable practical ongoing use, while conducting accuracy validation testing specific to current implementation would establish baseline performance metrics.

Enhanced AI integration improvements include transitioning to locally-hosted models like Llama for complete data privacy control. Implementing two-tier extraction with automatic quality checking and confidence scoring would help users and clinicians understand extraction reliability for individual data points.

Longer-term production readiness requires comprehensive HIPAA compliance implementation, integration APIs for existing Electronic Health Record systems, and multi-language support for diverse patient populations. Clinical validation studies with real patient cohorts would establish effectiveness metrics for patient engagement improvement, while provider dashboards could enable healthcare teams to monitor engagement patterns and data quality trends.

## Implementation & Ethics Memo: AI Health Data Engine

### How AI Assisted My Development Process

Building this health data extraction tool provided substantial insights into effective AI-assisted development practices. I relied primarily on Google AI Studio and Gemini throughout the development process, using AI both for the product's core functionality and for development assistance. The experience revealed important nuances in the relationship between AI capabilities and human oversight.

For code generation tasks, I discovered that specificity dramatically affected output quality. Broad requests like "create a health data dashboard" consistently yielded generic, unusable results. However, targeted prompts such as "generate a React component for voice input with SpeechRecognition API integration, including proper cleanup and error handling" produced solid foundational code requiring minimal modification. AI proved particularly effective at generating component structures, implementing specific browser APIs, and creating TypeScript interface definitions.

The most productive development approach involved using AI for discrete, well-defined implementation tasks while retaining human control over architectural decisions and component integration. AI consistently struggled with multi-step tasks requiring understanding of broader system context. When I requested styling changes alongside functional modifications, it often broke existing functionality while implementing visual updates. This pattern forced me to adopt a more methodical workflow: leverage AI for specific, isolated tasks, then personally handle integration to ensure system coherence.

Debugging experiences highlighted similar effectiveness patterns. AI could identify syntax errors and suggest fixes for isolated code blocks reliably, but struggled with complex state management issues or problems spanning multiple components. All design and user experience decisions remained entirely human-driven, from clinical aesthetic choices to information hierarchy and workflow design decisions.

The prompting evolution throughout development demonstrated that initial broad requests produced generic results, while specific, task-oriented prompts yielded practical code. This required developing skills in decomposing complex requirements into discrete, actionable AI tasks while maintaining human oversight of system architecture and user experience.

## Design Decisions and Implementation Choices

The core AI functionality - transforming unstructured patient narratives into structured JSON - directly addresses the fundamental friction point in health data collection identified through my thesis research. Patient disengagement stems primarily from cognitive burden imposed by structured data entry interfaces. Enabling natural language input eliminates this friction while providing clinicians with standardized data formats they require.

Several implementation decisions were made for practical reasons rather than ideal solutions. The most significant compromise involved eliminating backend infrastructure entirely. While production systems require secure, local processing to protect patient data, building complete backend architecture would have diverted focus from validating core AI extraction functionality. Direct frontend-to-API integration enables rapid testing and demonstration while acknowledging privacy limitations.

The tool calling methodology used for AI extraction comes directly from my thesis work, where I systematically tested and refined prompting strategies to maximize accuracy in clinical data extraction. Rather than rebuilding this validated prompting infrastructure, I leveraged established approaches, ensuring extraction accuracy consistent with previous research findings.

Including voice input alongside text entry reflects recognition that patients prefer different interaction modalities. Some individuals find speaking more natural than typing, particularly when describing symptoms or daily experiences. Browser SpeechRecognition API implementation added this accessibility without significantly complicating the AI processing pipeline.

The summary display format emerged from considering how clinicians actually consume patient-generated data. Rather than requiring users to parse raw JSON output, the categorized summary presents information in clinically relevant groupings that align with healthcare provider assessment patterns during patient consultations.

## Ethical Considerations and Risk Assessment

The current implementation presents several significant privacy and security concerns requiring transparent acknowledgment. Most critically, patient narratives transmit to Google's external API infrastructure. While data contains no explicit patient identifiers, health narratives inherently include sensitive personal information. This approach remains acceptable for research demonstration purposes but would be entirely inappropriate for clinical deployment.

I made this privacy trade-off consciously to enable rapid prototyping and concept validation. Production-ready systems would require either local AI model deployment or secure, HIPAA-compliant cloud infrastructure with proper encryption and access controls. The current architecture serves as proof-of-concept while clearly flagging fundamental changes necessary before clinical implementation.

Bias and fairness considerations present more subtle but potentially significant challenges. The example narrative generation feature may perpetuate model biases regarding how patients describe symptoms or what constitutes typical health experiences. Generated examples could inadvertently influence real patient reporting patterns in ways that skew data collection. Production systems would require example content curated from authentic, diverse patient voices rather than AI-generated content.

Over-reliance represents perhaps the most critical long-term concern. While the tool explicitly disclaims providing medical advice, risks remain that patients or clinical staff might trust AI extraction without appropriate validation. The structured JSON output appears authoritative and clean, potentially masking extraction errors or important omissions. Future implementations require confidence scoring, uncertainty indicators, and systematic validation workflows to prevent inappropriate reliance on AI outputs.

Regarding academic integrity in my development process, I maintained clear boundaries between AI assistance and original intellectual contribution. AI helped generate code scaffolding and implement specific technical requirements, but all conceptual decisions, user experience design, and integration work reflected my own analysis and judgment. The core AI methodology derives from my thesis research rather than AI-generated approaches, representing appropriate use of AI as productivity enhancement while ensuring fundamental problem-solving remained human-driven.

# Lessons Learned About Building with AI

The most significant discovery involved how task specificity dramatically affected AI output quality. Early project phases involved attempting complex, multi-faceted requests that consistently produced generic, poorly-integrated results. Learning to break requirements into discrete, specific tasks with clear inputs and outputs transformed AI effectiveness. This required developing new skills in translating conceptual understanding into specific, actionable prompts that AI could execute reliably.

Another unexpected finding was that AI excelled at implementing specific technical requirements but lacked design consistency and user experience coherence capabilities. While AI could generate individual components meeting functional specifications, it could not maintain aesthetic consistency across components or understand how design decisions affect overall user journey. This experience reinforced the critical importance of human oversight for ensuring cohesive product experiences.

For other founders considering AI tool integration, I would emphasize three key principles. First, invest upfront time in learning effective requirement decomposition into specific, isolated tasks that AI can handle reliably. Second, use AI for implementation and scaffolding while retaining human control over architectural decisions, user experience design, and component integration. Third, develop systematic validation processes since AI-generated code often works for tested cases but fails in edge conditions or when integrated with other systems.

This project fundamentally changed my perspective on AI's role in future ventures. Rather than viewing AI as replacement for technical expertise, I now understand it as powerful acceleration tool for implementation once conceptual frameworks and requirements are clearly defined through human analysis. For capstone work, this suggests focusing human effort on problem analysis, user research, and strategic decisions while leveraging AI for rapid implementation and technical iteration.

The experience reinforced that understanding both capabilities and limitations of available tools remains essential for effective use. AI assistance proves most valuable when you can clearly articulate requirements and evaluate whether outputs meet those specifications. This requires maintaining sufficient technical knowledge to guide and validate AI assistance rather than accepting outputs without critical assessment.

Most importantly, this project demonstrated that impactful AI applications may focus on eliminating tedious, repetitive tasks that create user friction rather than replacing human decision-making entirely. Using AI to bridge gaps between natural human communication preferences and system processing requirements can create more human-centered experiences while maintaining data quality necessary for effective outcomes. This approach feels more sustainable and ethically sound than attempting to automate human judgment processes entirely.